

lab04

June 27, 2025

1 Lab 04: Clustering and Association Rule Learning

1.0.1 E20280

1 Clustering

1.1 Importing Required Modules

```
[ ]: from sklearn.cluster import KMeans
      from sklearn.datasets import make_blobs # to generate sample datasets
```

1.2 Creating a Sample Dataset with 4 Clusters

```
[ ]: X, y = make_blobs (
      n_samples =400 ,
      centers =4 ,
      cluster_std =0.90 ,
      random_state =1
      )
```


1.3 Determining the Optimum Value of k Using the Elbow Method

```
[ ]: import matplotlib.pyplot as plt
      wcss = [] # within cluster sum of squares
      for i in range(1, 11):
          kmeans = KMeans(
              n_clusters=i,
              init="k-means++",
              max_iter=300,
              n_init=10,
              random_state=0
          )

          kmeans.fit(X)
          wcss.append(kmeans.inertia_)

      plt.plot(range(1, 11), wcss)
      plt.title("Elbow Method")
      plt.xlabel("Number of clusters")
```

```
plt.ylabel("WCSS")
plt.show()
```




1.4 Applying the K-Means Algorithm

```
[ ]: kmeans = KMeans(n_clusters=4, random_state=0) # from Elbow method
      closest_cluster_index = kmeans.fit_predict(X)
      cluster_centers = kmeans.cluster_centers_
```

1.5 Visualizing Clusters

```
[ ]: plt.scatter(X[:, 0], X[:, 1], c="green")
      plt.scatter(
          cluster_centers[:, 0],
          cluster_centers[:, 1],
          s=200,
          c="black",
          marker="*"
      )
      plt.show()
```



1.0.2 Exercise01

1. Import the iris dataset from scikit-learn. Convert it into an unlabeled dataset by removing the class attribute.

```
[ ]: from sklearn.datasets import load_iris

iris = load_iris()
iris_X = iris.data # features only, unlabeled dataset
```

2. Use the Elbow method to identify the best value for k (minimizing WCSS).

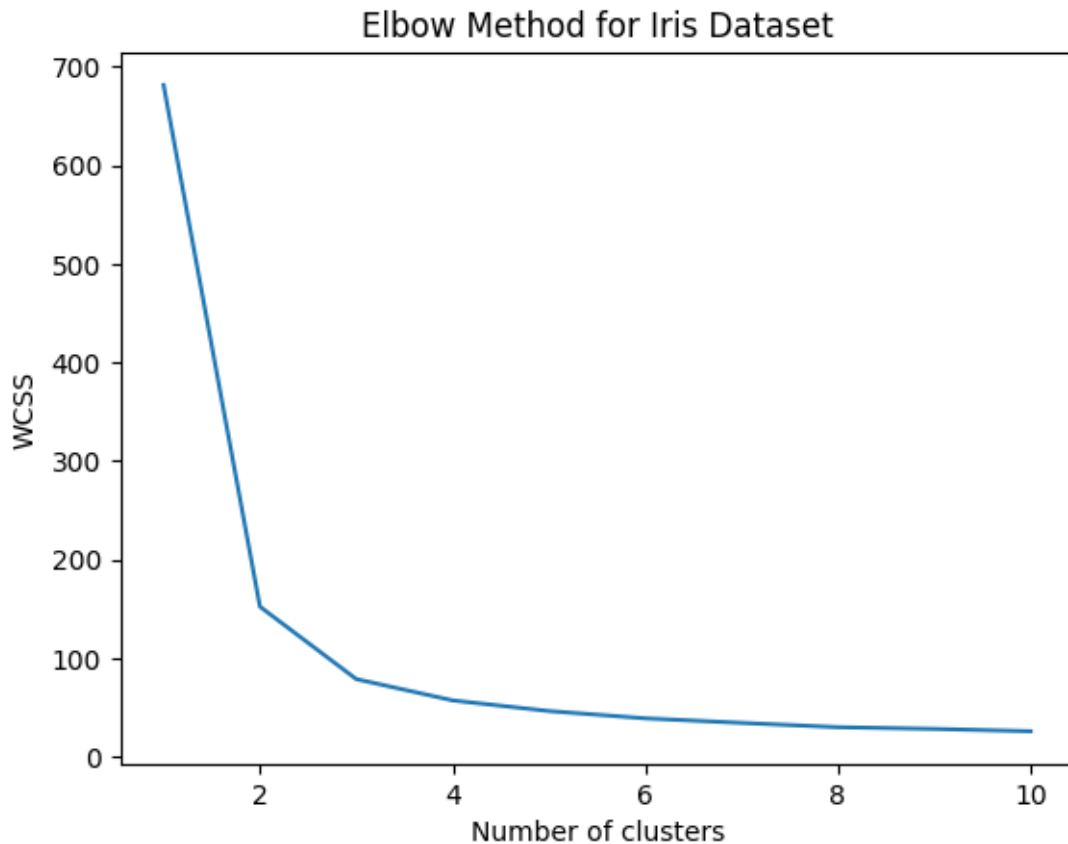
```
[ ]: wcss_iris = [] # Within Cluster Sum of Squares for iris dataset
# Using the Elbow Method to find the optimal number of clusters for the Iris_
↳ dataset
for k in range(1, 11):
    # kmeans with k clusters
    kmeans_iris = KMeans(n_clusters=k, init="k-means++", max_iter=300,
↳ n_init=10, random_state=0)
    # fit the model (train the model)
    kmeans_iris.fit(iris_X)
```

```

# append the inertia (WCSS) to the list (Inertia is the sum of squared
↪ distances of samples to their closest cluster center)
wcss_iris.append(kmeans_iris.inertia_)

plt.plot(range(1, 11), wcss_iris)
plt.title("Elbow Method for Iris Dataset")
plt.xlabel("Number of clusters")
plt.ylabel("WCSS")
plt.show()

```



3. Fit the K-Means algorithm with the k found in part (b)

```

[ ]: # Fit KMeans with the optimal k (3) found from the Elbow method for the iris
↪ dataset
kmeans_iris = KMeans(n_clusters=3, init="k-means++", max_iter=300, n_init=10,
↪ random_state=0)
iris_clusters = kmeans_iris.fit_predict(iris_X)

```

4. Explain the output of: 'kmeans . cluster_centers_'

The output of `kmeans.cluster_centers_` is a NumPy array containing the coordinates of the

centroids (centers) of each cluster found by the K-Means algorithm. Each row in this array represents the center of a cluster in the feature space, and each column corresponds to a feature. These centroids are the mean positions of all the points assigned to each cluster and are used to define the clusters in the data.

5. Visualize the data points and cluster centers in a 3D plot using the first three features as axes.


```
[ ]: from mpl_toolkits.mplot3d import Axes3D

fig = plt.figure(figsize=(12, 12))
ax = fig.add_subplot(111, projection='3d')

# Plot data points, colored by cluster assignment
ax.scatter(
    iris_X[:, 0], iris_X[:, 1], iris_X[:, 2],
    c=iris_clusters, cmap='viridis', s=40, alpha=0.7
)

# Plot cluster centers
centers = kmeans_iris.cluster_centers_
ax.scatter(
    centers[:, 0], centers[:, 1], centers[:, 2],
    c='red', s=200, marker='*', label='Centers'
)

ax.set_xlabel('Sepal length (cm)')
ax.set_ylabel('Sepal width (cm)')
ax.set_zlabel('Petal length (cm)')
ax.set_title('Iris Clusters (3D)')
ax.legend()
plt.show()
```



2 AssociationRuleLearning

2.1 InstallingApriori

```
[ ]: !pip install mlxtend
```

Requirement already satisfied: mlxtend in e:\my projects\python\co544 machine learning\venv\lib\site-packages (0.23.4)

Requirement already satisfied: pandas>=0.24.2 in e:\my projects\python\co544 machine learning\venv\lib\site-packages (from mlxtend) (2.2.3)

Requirement already satisfied: joblib>=0.13.2 in e:\my projects\python\co544 machine learning\venv\lib\site-packages (from mlxtend) (1.4.2)

Requirement already satisfied: numpy>=1.16.2 in e:\my projects\python\co544

machine learning\venv\lib\site-packages (from mlxtend) (2.0.2)
 Requirement already satisfied: matplotlib>=3.0.0 in e:\my projects\python\co544
 machine learning\venv\lib\site-packages (from mlxtend) (3.9.4)
 Requirement already satisfied: scikit-learn>=1.3.1 in e:\my
 projects\python\co544 machine learning\venv\lib\site-packages (from mlxtend)
 (1.6.1)
 Requirement already satisfied: scipy>=1.2.1 in e:\my projects\python\co544
 machine learning\venv\lib\site-packages (from mlxtend) (1.13.1)
 Requirement already satisfied: fonttools>=4.22.0 in e:\my projects\python\co544
 machine learning\venv\lib\site-packages (from matplotlib>=3.0.0->mlxtend)
 (4.57.0)
 Requirement already satisfied: cycler>=0.10 in e:\my projects\python\co544
 machine learning\venv\lib\site-packages (from matplotlib>=3.0.0->mlxtend)
 (0.12.1)
 Requirement already satisfied: importlib-resources>=3.2.0 in e:\my
 projects\python\co544 machine learning\venv\lib\site-packages (from
 matplotlib>=3.0.0->mlxtend) (6.5.2)
 Requirement already satisfied: pyparsing>=2.3.1 in e:\my projects\python\co544
 machine learning\venv\lib\site-packages (from matplotlib>=3.0.0->mlxtend)
 (3.2.3)
 Requirement already satisfied: kiwisolver>=1.3.1 in e:\my projects\python\co544
 machine learning\venv\lib\site-packages (from matplotlib>=3.0.0->mlxtend)
 (1.4.7)
 Requirement already satisfied: packaging>=20.0 in e:\my projects\python\co544
 machine learning\venv\lib\site-packages (from matplotlib>=3.0.0->mlxtend)
 (25.0)
 Requirement already satisfied: python-dateutil>=2.7 in e:\my
 projects\python\co544 machine learning\venv\lib\site-packages (from
 matplotlib>=3.0.0->mlxtend) (2.9.0.post0)
 Requirement already satisfied: pillow>=8 in e:\my projects\python\co544 machine
 learning\venv\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (11.2.1)
 Requirement already satisfied: contourpy>=1.0.1 in e:\my projects\python\co544
 machine learning\venv\lib\site-packages (from matplotlib>=3.0.0->mlxtend)
 (1.3.0)
 Requirement already satisfied: pytz>=2020.1 in e:\my projects\python\co544
 machine learning\venv\lib\site-packages (from pandas>=0.24.2->mlxtend) (2025.2)
 Requirement already satisfied: tzdata>=2022.7 in e:\my projects\python\co544
 machine learning\venv\lib\site-packages (from pandas>=0.24.2->mlxtend) (2025.2)
 Requirement already satisfied: threadpoolctl>=3.1.0 in e:\my
 projects\python\co544 machine learning\venv\lib\site-packages (from scikit-
 learn>=1.3.1->mlxtend) (3.6.0)
 Requirement already satisfied: zipp>=3.1.0 in e:\my projects\python\co544
 machine learning\venv\lib\site-packages (from importlib-
 resources>=3.2.0->matplotlib>=3.0.0->mlxtend) (3.21.0)
 Requirement already satisfied: six>=1.5 in e:\my projects\python\co544 machine
 learning\venv\lib\site-packages (from python-
 dateutil>=2.7->matplotlib>=3.0.0->mlxtend) (1.17.0)

WARNING: You are using pip version 22.0.4; however, version 25.1.1 is available.
 You should consider upgrading via the 'E:\My Projects\python\C0544 Machine Learning\venv\Scripts\python.exe -m pip install --upgrade pip' command.

2.2 Importing Required Modules

```
[ ]: from mlxtend . frequent_patterns import apriori
      from mlxtend . frequent_patterns import association_rules
      from mlxtend . preprocessing import TransactionEncoder
      import pandas as pd
```

2.3 Input Data

```
[ ]: dataset = [
      ["Milk", "Onion", "Nutmeg", "Kidney Beans", "Eggs", "Yogurt"],
      ["Dill", "Onion", "Nutmeg", "Kidney Beans", "Eggs", "Yogurt"],
      ["Milk", "Apple", "Kidney Beans", "Eggs"],
      ["Milk", "Unicorn", "Corn", "Kidney Beans", "Yogurt"],
      ["Corn", "Onion", "Onion", "Kidney Beans", "Ice cream", "Eggs"]
    ]
```

2.4 Creating the DataFrame of Frequent Itemsets

```
[ ]: te = TransactionEncoder()
      te_ary = te.fit( dataset ).transform( dataset )
      df = pd.DataFrame( te_ary , columns = te . columns_ )
```

2.5 Applying Apriori Algorithm and Finding Association Rules

```
[ ]: freq = apriori( df , min_support =0.002 , use_colnames = True )
      rules = association_rules( freq , metric ="lift" , min_threshold =1)
```

```
e:\My Projects\python\C0544 Machine Learning\venv\lib\site-
packages\mlxtend\frequent_patterns\association_rules.py:186: RuntimeWarning:
invalid value encountered in divide
      cert_metric = np.where(certainty_denom == 0, 0, certainty_num /
certainty_denom)
```

1.0.3 Exercise02

1. Import the provided groceries.csv dataset.

```
[ ]: groceries_df = pd.read_csv('groceries.csv')
      groceries_df.head()
```

```
[ ]:
      citrus fruit semi-finished bread      margarine \
0      tropical fruit      yogurt      coffee
1      whole milk      NaN      NaN
2      pip fruit      yogurt      cream cheese
3      other vegetables      whole milk      condensed milk
```


4	whole milk		butter		yogurt	
		ready soups		Unnamed: 4	Unnamed: 5	Unnamed: 6 \
0		NaN		NaN	NaN	NaN
1		NaN		NaN	NaN	NaN
2		meat spreads		NaN	NaN	NaN
3	long life bakery product			NaN	NaN	NaN
4		rice	abrasive cleaner		NaN	NaN
	Unnamed: 7	Unnamed: 8	Unnamed: 9 ...	Unnamed: 22	Unnamed: 23	Unnamed: 24 \
0	NaN	NaN	NaN ...	NaN	NaN	NaN
1	NaN	NaN	NaN ...	NaN	NaN	NaN
2	NaN	NaN	NaN ...	NaN	NaN	NaN
3	NaN	NaN	NaN ...	NaN	NaN	NaN
4	NaN	NaN	NaN ...	NaN	NaN	NaN
	Unnamed: 25	Unnamed: 26	Unnamed: 27	Unnamed: 28	Unnamed: 29	Unnamed: 30 \
0	NaN	NaN	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN	NaN	NaN
	Unnamed: 31					
0	NaN					
1	NaN					
2	NaN					
3	NaN					
4	NaN					

[5 rows x 32 columns]

2. Explore the dataset and build the frequent-item DataFrame.

```
[ ]: # Explore the groceries_df dataset
print("Number of transactions:", len(groceries_df))
print("Sample transactions:")
print(groceries_df.head())

# Convert the DataFrame into a list of transactions (dropping NaNs)
transactions = groceries_df.apply(lambda row: [item for item in row if pd.
    notnull(item)], axis=1).tolist()

# Use TransactionEncoder to encode the transactions
te_groceries = TransactionEncoder()
te_groceries_ary = te_groceries.fit(transactions).transform(transactions)
groceries_onehot = pd.DataFrame(te_groceries_ary, columns=te_groceries.columns_)
```

```
# Build the frequent-item DataFrame using apriori
freq_groceries = apriori(groceries_onehot, min_support=0.01, use_colnames=True)
freq_groceries.head()
```

Number of transactions: 9834

Sample transactions:

	citrus fruit	semi-finished bread		margarine	\			
0	tropical fruit	yogurt		coffee				
1	whole milk	NaN		NaN				
2	pip fruit	yogurt		cream cheese				
3	other vegetables	whole milk		condensed milk				
4	whole milk	butter		yogurt				
	ready soups	Unnamed: 4	Unnamed: 5	Unnamed: 6	\			
0	NaN	NaN	NaN	NaN				
1	NaN	NaN	NaN	NaN				
2	meat spreads	NaN	NaN	NaN				
3	long life bakery product	NaN	NaN	NaN				
4	rice	abrasive cleaner	NaN	NaN				
	Unnamed: 7	Unnamed: 8	Unnamed: 9	...	Unnamed: 22	Unnamed: 23	Unnamed: 24	\
0	NaN	NaN	NaN	...	NaN	NaN	NaN	
1	NaN	NaN	NaN	...	NaN	NaN	NaN	
2	NaN	NaN	NaN	...	NaN	NaN	NaN	
3	NaN	NaN	NaN	...	NaN	NaN	NaN	
4	NaN	NaN	NaN	...	NaN	NaN	NaN	
	Unnamed: 25	Unnamed: 26	Unnamed: 27	Unnamed: 28	Unnamed: 29	Unnamed: 30	\	
0	NaN	NaN	NaN	NaN	NaN	NaN		
1	NaN	NaN	NaN	NaN	NaN	NaN		
2	NaN	NaN	NaN	NaN	NaN	NaN		
3	NaN	NaN	NaN	NaN	NaN	NaN		
4	NaN	NaN	NaN	NaN	NaN	NaN		
	Unnamed: 31							
0	NaN							
1	NaN							
2	NaN							
3	NaN							
4	NaN							

[5 rows x 32 columns]

```
[ ]:      support      itemsets
0  0.033455      (UHT-milk)
1  0.017694  (baking powder)
```

```

2  0.052471      (beef)
3  0.033252      (berries)
4  0.026032      (beverages)

```

3. Apply the Apriori algorithm to find itemsets with support > 8%.

```

[ ]: # Find frequent itemsets in groceries_onehot with support greater than 8%
freq_groceries_08 = apriori(groceries_onehot, min_support=0.08,
↪use_colnames=True)
freq_groceries_08

```

```

[ ]:      support      itemsets
0  0.080537      (bottled beer)
1  0.110535      (bottled water)
2  0.082672      (citrus fruit)
3  0.193512 (other vegetables)
4  0.088977      (pastry)
5  0.183954      (rolls/buns)
6  0.109010      (root vegetables)
7  0.093960      (sausage)
8  0.098536      (shopping bags)
9  0.174395      (soda)
10 0.104942      (tropical fruit)
11 0.255542      (whole milk)
12 0.139516      (yogurt)

```

4. Generate association rules using the lift metric.

```

[ ]: # Generate association rules for the groceries dataset using the lift metric
rules_groceries = association_rules(freq_groceries, metric="lift",
↪min_threshold=1)
rules_groceries.head()

```

```

[ ]:      antecedents      consequents  antecedent support \
0      (beef) (other vegetables)      0.052471
1 (other vegetables)      (beef)      0.193512
2      (beef)      (rolls/buns)      0.052471
3 (rolls/buns)      (beef)      0.183954
4      (beef)      (root vegetables)      0.052471

      consequent support      support      confidence      lift      representativity \
0      0.193512 0.019727      0.375969      1.942869      1.0
1      0.052471 0.019727      0.101944      1.942869      1.0
2      0.183954 0.013626      0.259690      1.411714      1.0
3      0.052471 0.013626      0.074074      1.411714      1.0
4      0.109010 0.017389      0.331395      3.040058      1.0

```

	leverage	conviction	zhangs_metric	jaccard	certainty	kulczynski
0	0.009574	1.292384	0.512171	0.087191	0.226236	0.238957
1	0.009574	1.055089	0.601742	0.087191	0.052213	0.238957
2	0.003974	1.102303	0.307791	0.061159	0.092809	0.166882
3	0.003974	1.023331	0.357383	0.061159	0.022799	0.166882
4	0.011669	1.332612	0.708220	0.120677	0.249594	0.245455

5. Select one rule and interpret it in your own words.

```
[ ]: # Select the first rule from rules_groceries and display it
selected_rule = rules_groceries.iloc[0]
print("Selected Rule:")
print(selected_rule)

# Interpretation
print("\nInterpretation:")
print(f"If a customer buys {list(selected_rule['antecedents'])}, they are
↳likely to also buy {list(selected_rule['consequents'])}.")
print(f"This rule has a confidence of {selected_rule['confidence']:.2f} and a
↳lift of {selected_rule['lift']:.2f}, meaning the likelihood of buying
↳{list(selected_rule['consequents'])} increases by a factor of
↳{selected_rule['lift']:.2f} when {list(selected_rule['antecedents'])} is
↳purchased.")
```

Selected Rule:

antecedents	(beef)
consequents	(other vegetables)
antecedent support	0.052471
consequent support	0.193512
support	0.019727
confidence	0.375969
lift	1.942869
representativity	1.0
leverage	0.009574
conviction	1.292384
zhangs_metric	0.512171
jaccard	0.087191
certainty	0.226236
kulczynski	0.238957

Name: 0, dtype: object

Interpretation:

If a customer buys ['beef'], they are likely to also buy ['other vegetables']. This rule has a confidence of 0.38 and a lift of 1.94, meaning the likelihood of buying ['other vegetables'] increases by a factor of 1.94 when ['beef'] is purchased.

6. How many rules satisfy both lift > 4 and confidence > 0.8?

```
[ ]: # Count the number of rules with lift > 4 and confidence > 0.8 in the 'rules' DataFrame
num_rules = rules[(rules['lift'] > 4) & (rules['confidence'] > 0.8)].shape[0]
print("Number of rules with lift > 4 and confidence > 0.8:", num_rules)
```

Number of rules with lift > 4 and confidence > 0.8: 60